




A Study in Reduction


Parallel Reduction

- Common and important data parallel primitive
- Many data elements -> single output (associative!)
- Easy to implement in CUDA
- Tree-based approach



Parallel Reduction

- Need to use multiple thread blocks
 - To process large arrays
 - To fully utilize GPU
- Partition the array, one block per partition
- How to communicate partial results?



Reduction #1: AtomicGlobal

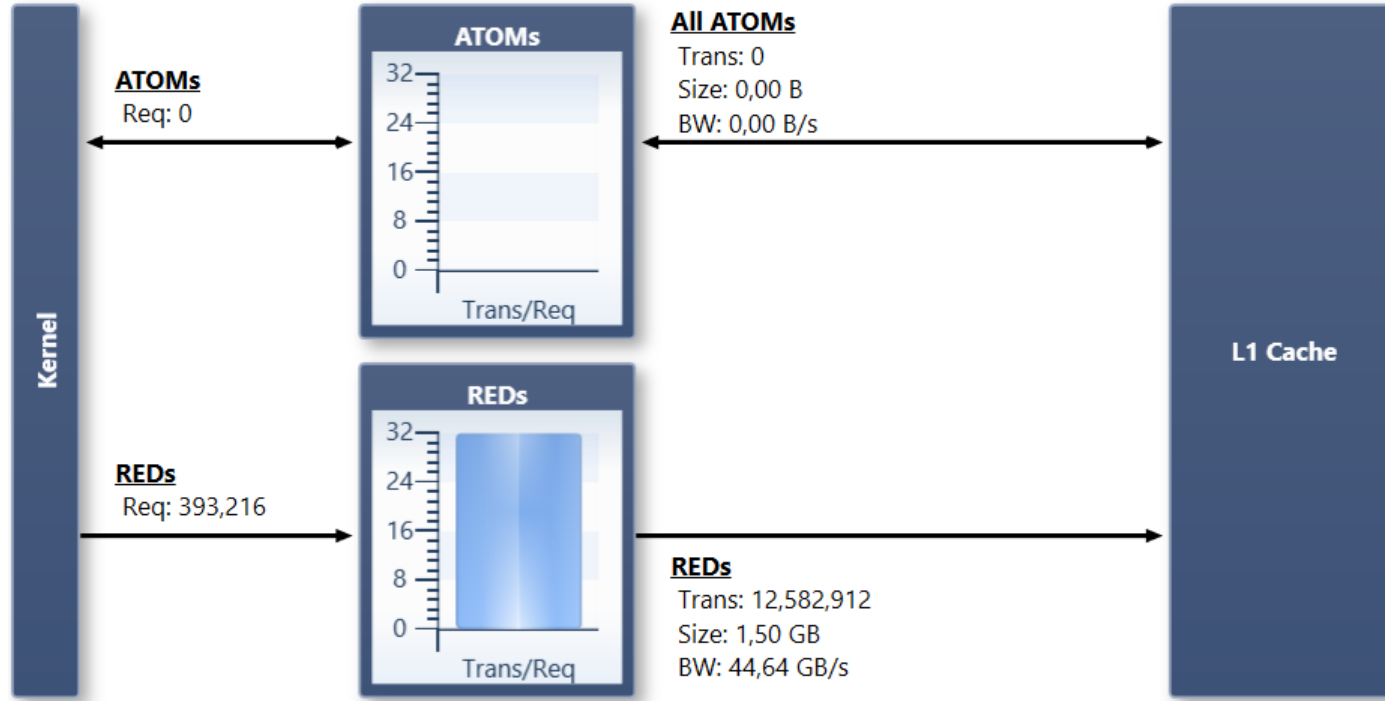
```
__global__ void reduceAtomicGlobal(const float* input,  
                                   float* result, int elements)  
{  
    int id = threadIdx.x + blockIdx.x*blockDim.x;  
    atomicAdd(result, input[id]);  
}
```

Performance 12M Reduction

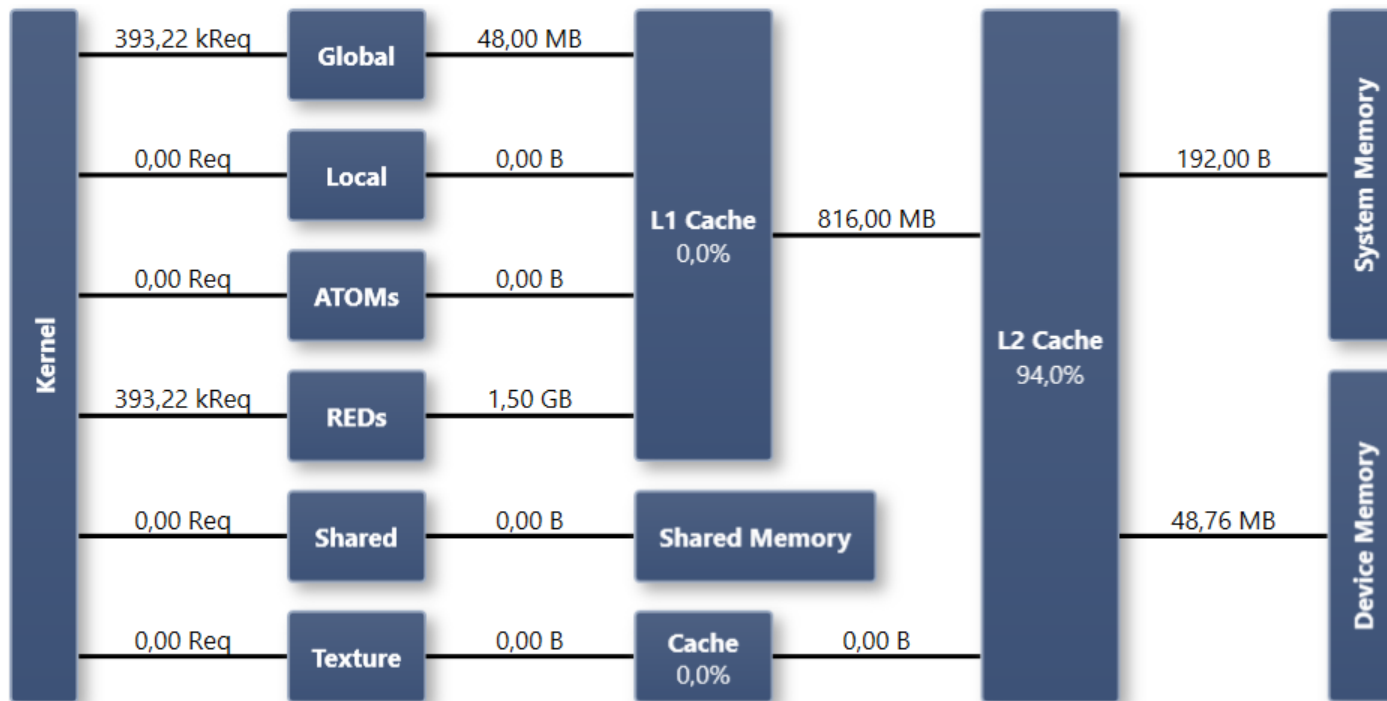
	Time (3*2 ²² floats)	Bandwidth	Step speedup	Cumulative Speedup	Speedup to reduction
AtomicGlobal	37.53ms	1.248GB/s	1.000	1.000	

GTX 680 peak: 192GB/S

Reduction #1: AtomicGlobal



Reduction #1: AtomicGlobal



Reduction #2: AtomicShared

```
__global__ void reduceAtomicShared(const float* input,
                                   float* result, int elements)
{
    int id = threadIdx.x + blockIdx.x*blockDim.x;
    float in = input[id];

    __shared__ float x;
    x = 0.0f;
    __syncthreads();

    atomicAdd(&x, in);
    __syncthreads();

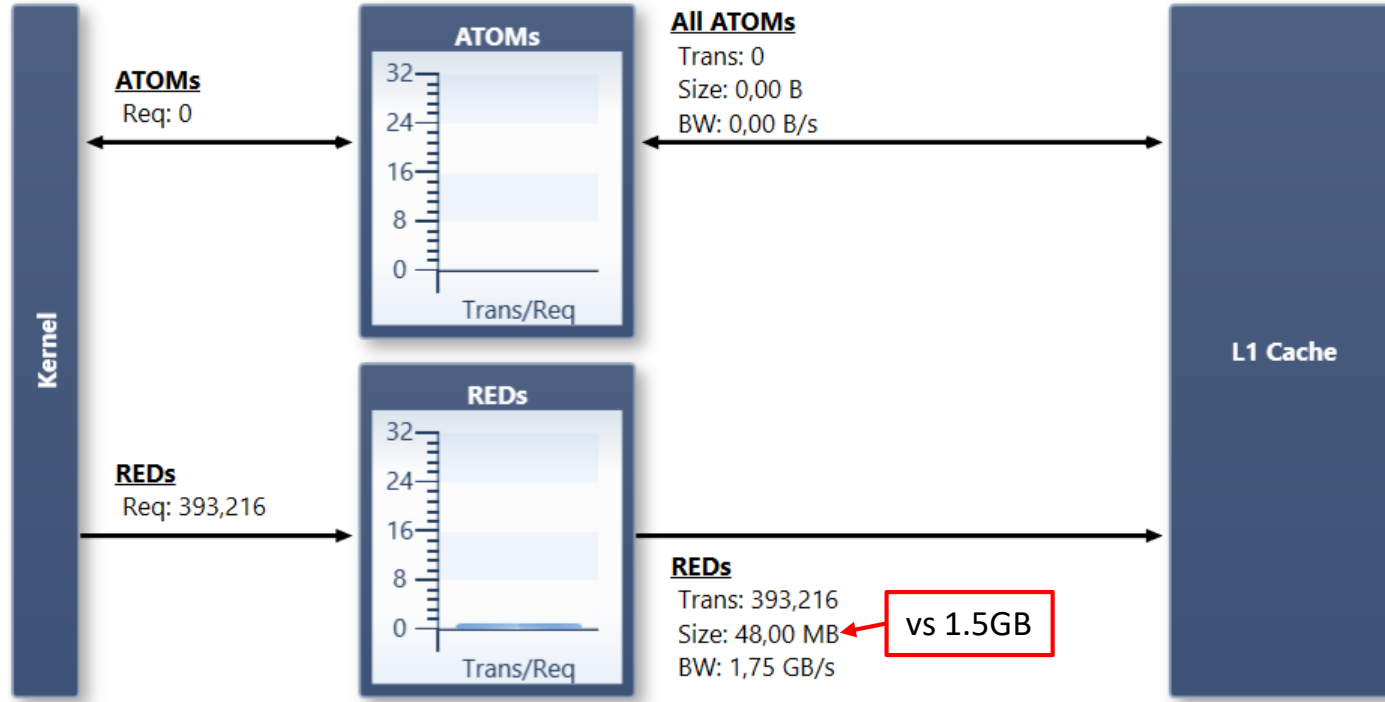
    if(threadIdx.x == 0)
        atomicAdd(result, x);
}
```


Performance 12M Reduction

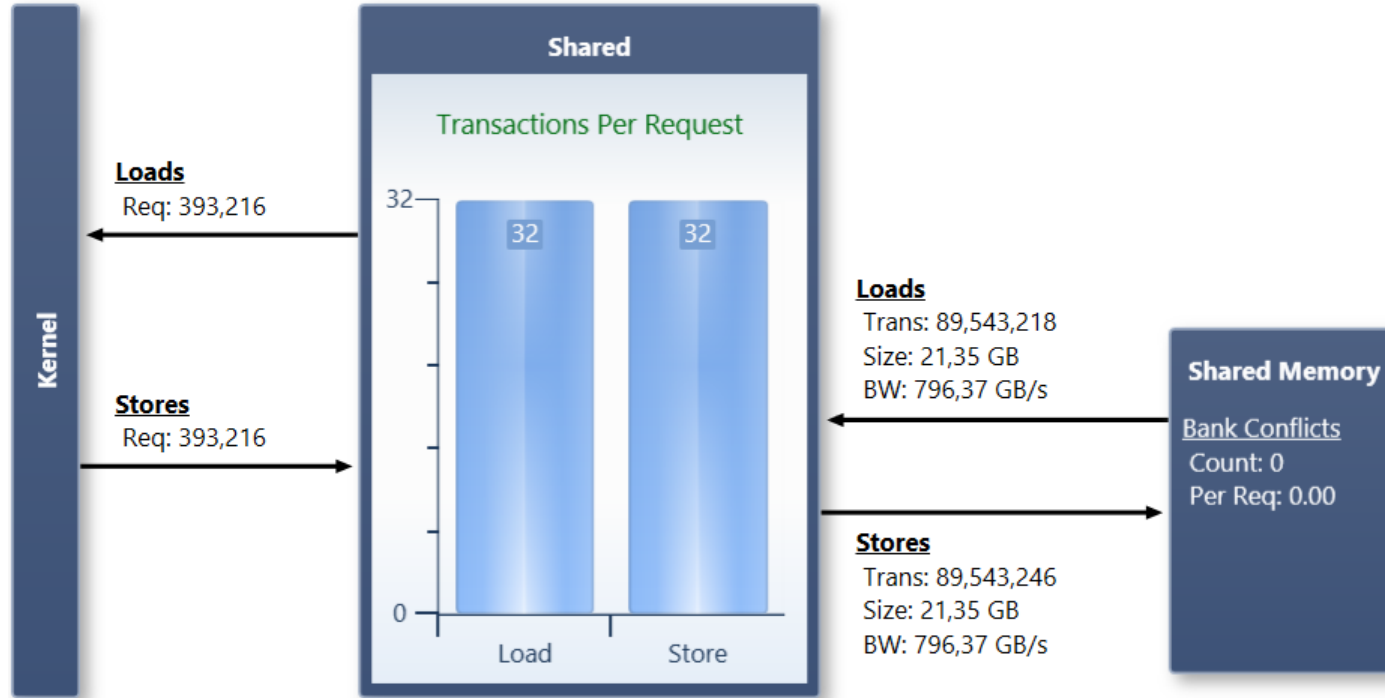
	Time ($3 \cdot 2^{22}$ floats)	Bandwidth	Step speedup	Cumulative Speedup	Speedup to reduction
AtomicGlobal	37.53ms	1.248GB/s	1.000	1.000	
AtomicShared	29.94ms	1.565GB/s	1.254	1.254	

GTX 680 peak: 192GB/S

Reduction #2: AtomicShared



Reduction #2: AtomicShared




Reduction #3: Shared

```
__global__ void reduceShared(const float* input,
                             float* result, int elements)
{
    extern __shared__ float data[];
    int id = threadIdx.x + blockIdx.x*blockDim.x;
    data[threadIdx.x] = input[id];
    __syncthreads();

    for(int s = 1; s < blockDim.x; s*=2)
    {
        if(threadIdx.x % (2*s) == 0)
            data[threadIdx.x] += data[threadIdx.x + s];
        __syncthreads();
    }
    if(threadIdx.x == 0)
        atomicAdd(result, x);
}
```

Reduction #3: Shared

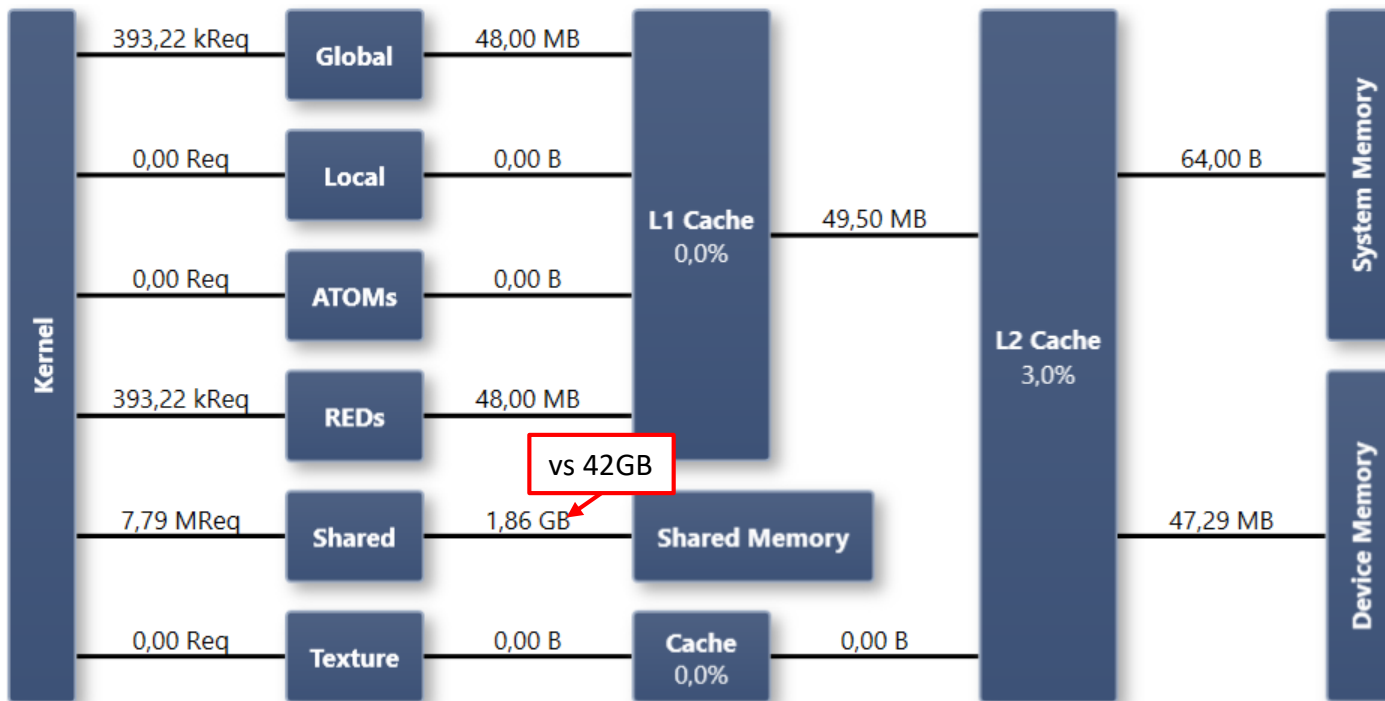


Performance 12M Reduction

	Time ($3 \cdot 2^{22}$ floats)	Bandwidth	Step speedup	Cumulative Speedup	Speedup to reduction
AtomicGlobal	37.53ms	1.248GB/s	1.000	1.000	
AtomicShared	29.94ms	1.565GB/s	1.254	1.254	
Reduction	4.452ms	10.52GB/s	6.725	8.43	1.000

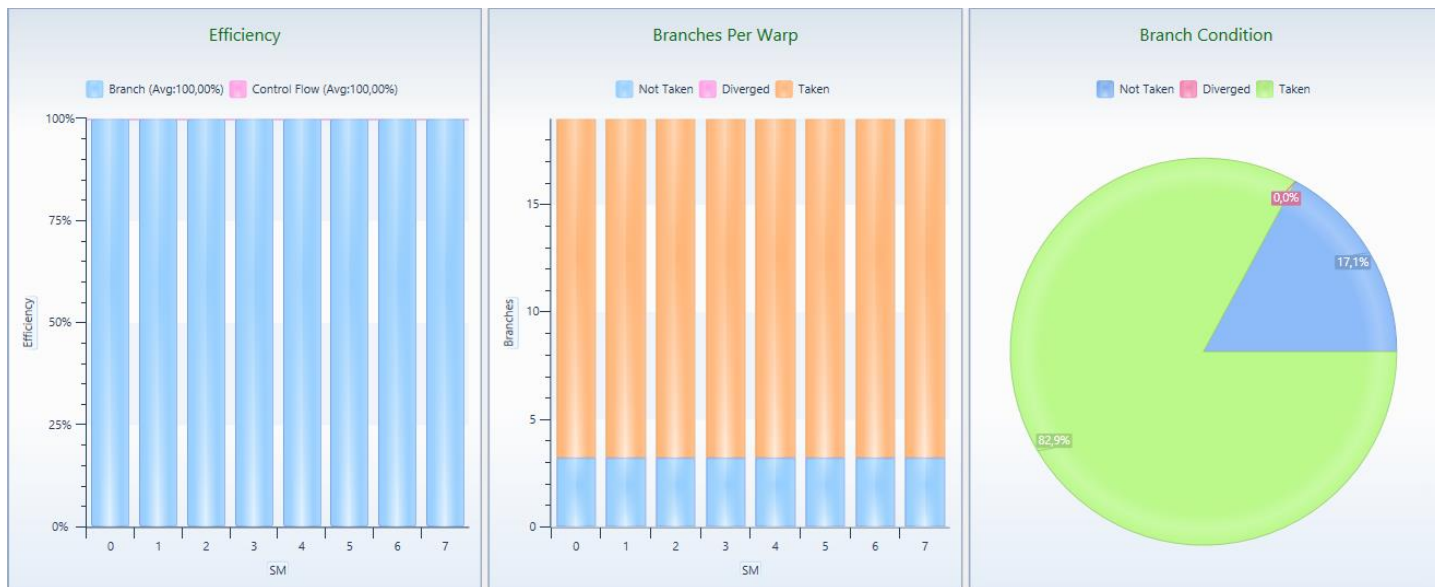
GTX 680 peak: 192GB/S

Reduction #3: Shared



Reduction #3: Shared

- Problem: Divergence




Reduction #4: SharedAntiDivergence

```
for(int s = 1; s < blockDim.x; s*=2)
{
    if(threadIdx.x % (2*s) == 0)
        data[threadIdx.x] += data[threadIdx.x + s];
    __syncthreads();
}
```



```
for(int s = 1; s < blockDim.x; s*=2)
{
    int lid = 2*s*threadIdx.x;
    if(lid < blockDim.x)
        data[lid] += data[lid + s];
    __syncthreads();
}
```

Reduction #4: SharedAntiDivergence





Performance 12M Reduction

	Time ($3 \cdot 2^{22}$ floats)	Bandwidth	Step speedup	Cumulative Speedup	Speedup to reduction
AtomicGlobal	37.53ms	1.248GB/s	1.000	1.000	
AtomicShared	29.94ms	1.565GB/s	1.254	1.254	
Reduction	4.452ms	10.52GB/s	6.725	8.43	1.000
AntiDivergence	3.604ms	13.00GB/s	1.235	10.41	1.235

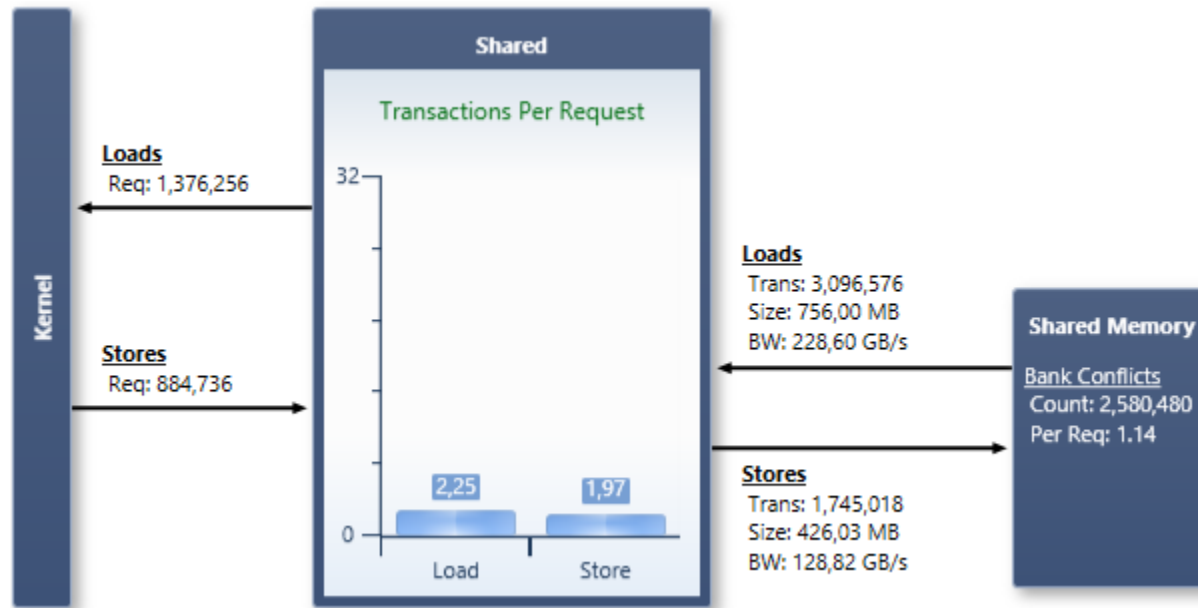
GTX 680 peak: 192GB/S

Reduction #4: SharedAntiDivergence



Reduction #4: SharedAntiDivergence

- Problem: bank conflicts




Reduction #5: SharedAntiConflicts

```
for(int s = 1; s < blockDim.x; s*=2)
{
    int lid = 2*s*threadIdx.x;
    if(lid < blockDim.x)
        data[lid] += data[lid + s];
    __syncthreads();
}
```



```
for(int s = blockDim.x/2; s > 0; s/=2)
{
    if(threadIdx.x < s)
        data[threadIdx.x] += data[threadIdx.x + s];
    __syncthreads();
}
```


Reduction #5: SharedAntiConflicts

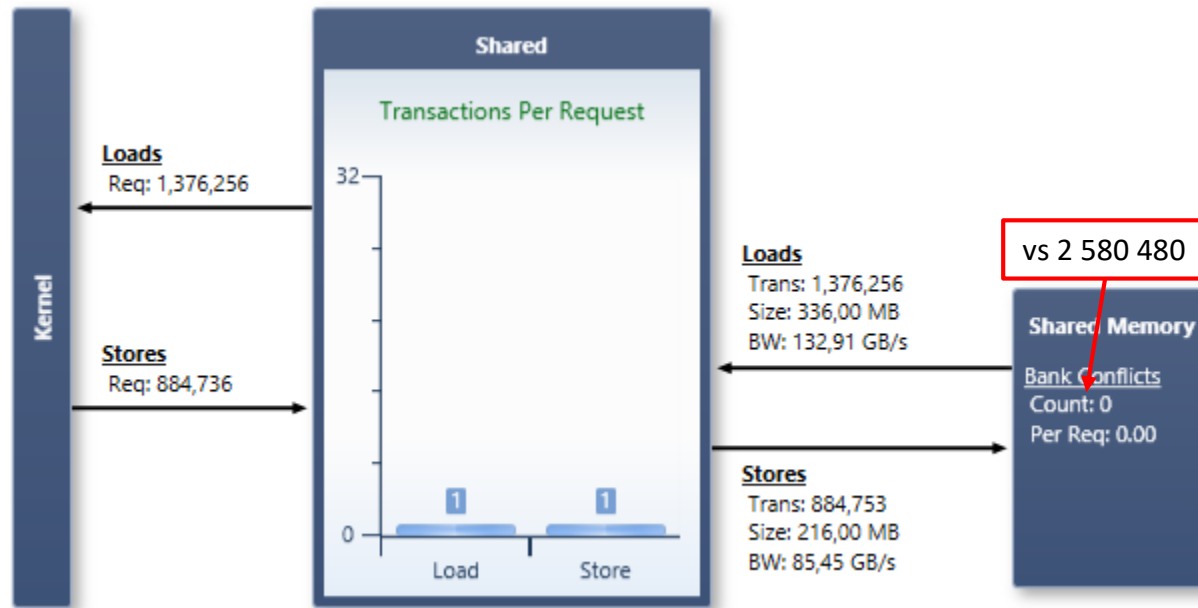


Performance 12M Reduction

	Time ($3 \cdot 2^{22}$ floats)	Bandwidth	Step speedup	Cumulative Speedup	Speedup to reduction
AtomicGlobal	37.53ms	1.248GB/s	1.000	1.000	
AtomicShared	29.94ms	1.565GB/s	1.254	1.254	
Reduction	4.452ms	10.52GB/s	6.725	8.43	1.000
AntiDivergence	3.604ms	13.00GB/s	1.235	10.41	1.235
AntiConflicts	2.751ms	17.03GB/s	1.310	13.64	1.618

GTX 680 peak: 192GB/S

Reduction #5: SharedAntiConflicts



Reduction #5: SharedAntiConflicts

- Problem: half of threads idle

```
for(int s = blockDim.x/2; s > 0; s/=2)
{
    if(threadIdx.x < s)
        data[threadIdx.x] += data[threadIdx.x + s];
    __syncthreads();
}
```

Reduction #6: SharedDualLoad

```
__global__ void reduceSharedDualLoad(const float* input,
                                     float* result, int elements)
{
    extern __shared__ float data[];
    int id = threadIdx.x + blockIdx.x*blockDim.x*2;
    data[threadIdx.x] = input[id] + input[id+blockDim.x];
    __syncthreads();
    for(int s = blockDim.x/2; s > 0; s/=2)
    {
        if(threadIdx.x < s)
            data[threadIdx.x] += data[threadIdx.x + s];
        __syncthreads();
    }
    if(threadIdx.x == 0)
        atomicAdd(result, data[0]);
}
```


Performance 12M Reduction

	Time ($3 \cdot 2^{22}$ floats)	Bandwidth	Step speedup	Cumulative Speedup	Speedup to reduction
AtomicGlobal	37.53ms	1.248GB/s	1.000	1.000	
AtomicShared	29.94ms	1.565GB/s	1.254	1.254	
Reduction	4.452ms	10.52GB/s	6.725	8.43	1.000
AntiDivergence	3.604ms	13.00GB/s	1.235	10.41	1.235
AntiConflicts	2.751ms	17.03GB/s	1.310	13.64	1.618
DualLoad	1.411ms	33.20GB/s	1.949	26.59	3.154


GTX 680 peak: 192GB/S

Reduction #6: SharedDualLoad

SharedAntiConflicts<<<24576,512>>>(...)



SharedDualLoad<<<12288,512>>>(...)



Reduction #6: SharedDualLoad

- 32GB/s vs 192GB/s peak
- What is the bottleneck?
 - Probably instruction overhead
 - Operations that are not loads, stores, or computations
 - Address arithmetic, loops, conditionals
- Idea: unroll loops

Reduction #7: SharedUnrolledLast

```
extern __shared__ volatile float data[];
...
for(int s = blockDim.x/2; s > 32; s/=2)
{
    if(threadIdx.x < s)
        data[threadIdx.x] += data[threadIdx.x + s];
    __syncthreads();
}

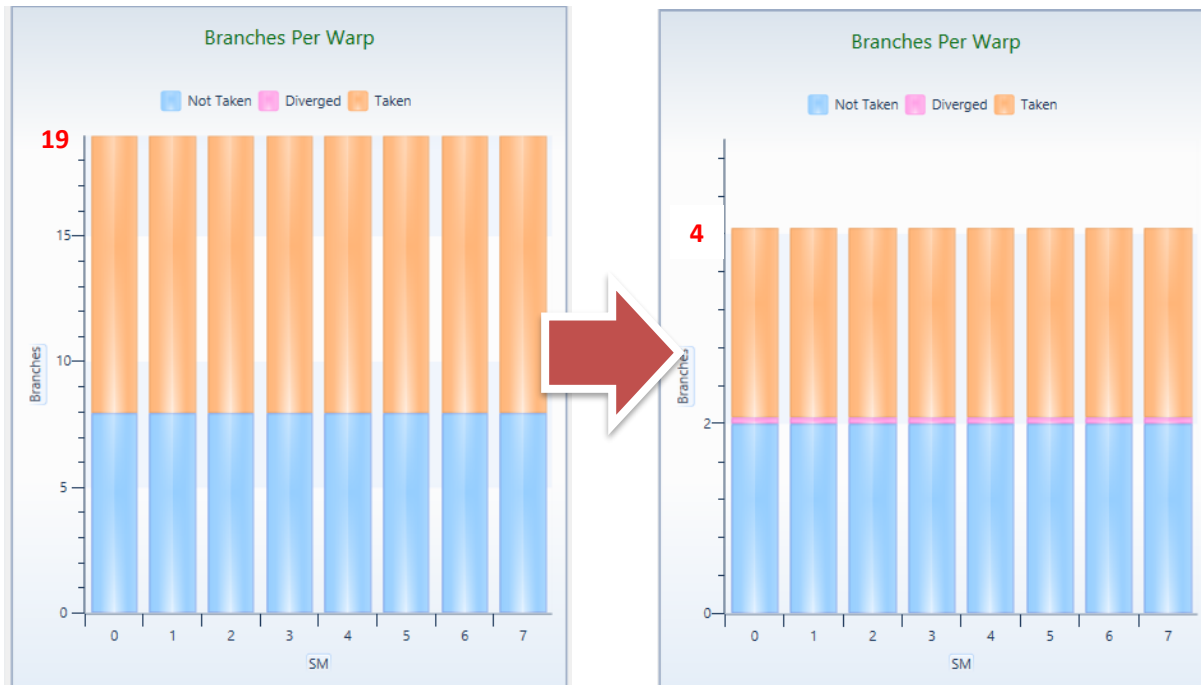
if(threadIdx.x < 32)
{
    data[threadIdx.x] += data[threadIdx.x + 32];
    data[threadIdx.x] += data[threadIdx.x + 16];
    data[threadIdx.x] += data[threadIdx.x + 8];
    data[threadIdx.x] += data[threadIdx.x + 4];
    data[threadIdx.x] += data[threadIdx.x + 2];
}
if(threadIdx.x == 0)
    atomicAdd(result, data[0] + data[1]);
```

Performance 12M Reduction

	Time ($3 \cdot 2^{22}$ floats)	Bandwidth	Step speedup	Cumulative Speedup	Speedup to reduction
AtomicGlobal	37.53ms	1.248GB/s	1.000	1.000	
AtomicShared	29.94ms	1.565GB/s	1.254	1.254	
Reduction	4.452ms	10.52GB/s	6.725	8.43	1.000
AntiDivergence	3.604ms	13.00GB/s	1.235	10.41	1.235
AntiConflicts	2.751ms	17.03GB/s	1.310	13.64	1.618
DualLoad	1.411ms	33.20GB/s	1.949	26.59	3.154
UnrolledLast	0.978ms	47.90GB/s	1.443	38.35	4.550

GTX 680 peak: 192GB/S

Reduction #7: SharedUnrolledLast



Reduction #8: SharedUnrolledAll (1/2)

```

template<int BlockSize>
__global__ void reduceSharedUnrolledAll(const float* input,
                                       float* result, int elements)
{
    extern __shared__ volatile float data[];
    int id = threadIdx.x + blockIdx.x*blockDim.x*2;
    data[threadIdx.x] = input[id] + input[id+blockDim.x];
    __syncthreads();

    if(BlockSize >= 512)
    {
        if(threadIdx.x < 256)
            data[threadIdx.x] += data[threadIdx.x + 256];
        __syncthreads();
    }
    if(BlockSize >= 256)
    {
        if(threadIdx.x < 128)
            data[threadIdx.x] += data[threadIdx.x + 128];
        __syncthreads();
    }
}

```

Reduction #8: SharedUnrolledAll (1/2)

```
if(BlockSize >= 128)
{
    if(threadIdx.x < 64)
        data[threadIdx.x] += data[threadIdx.x + 64];
    __syncthreads();
}

if(threadIdx.x < 32)
{
    data[threadIdx.x] += data[threadIdx.x + 32];
    data[threadIdx.x] += data[threadIdx.x + 16];
    data[threadIdx.x] += data[threadIdx.x + 8];
    data[threadIdx.x] += data[threadIdx.x + 4];
    data[threadIdx.x] += data[threadIdx.x + 2];
}


if(threadIdx.x == 0)
    atomicAdd(result, data[0] + data[1]);
}
```

Performance 12M Reduction

	Time ($3 \cdot 2^{22}$ floats)	Bandwidth	Step speedup	Cumulative Speedup	Speedup to reduction
AtomicGlobal	37.53ms	1.248GB/s	1.000	1.000	
AtomicShared	29.94ms	1.565GB/s	1.254	1.254	
Reduction	4.452ms	10.52GB/s	6.725	8.43	1.000
AntiDivergence	3.604ms	13.00GB/s	1.235	10.41	1.235
AntiConflicts	2.751ms	17.03GB/s	1.310	13.64	1.618
DualLoad	1.411ms	33.20GB/s	1.949	26.59	3.154
UnrolledLast	0.978ms	47.90GB/s	1.443	38.35	4.550
UnrolledAll	0.821ms	57.02GB/s	1.191	45.66	5.417

GTX 680 peak: 192GB/S

Reduction #8: SharedUnrolledAll



Reduction #9: SharedShuffle

```
if(threadIdx.x < 32)
{
    data[threadIdx.x] += data[threadIdx.x + 32];
    data[threadIdx.x] += data[threadIdx.x + 16];
    data[threadIdx.x] += data[threadIdx.x + 8];
    data[threadIdx.x] += data[threadIdx.x + 4];
    data[threadIdx.x] += data[threadIdx.x + 2];
}
```



```
if(threadIdx.x < 32)
{
    d = data[threadIdx.x] + data[threadIdx.x + 32];
    d += __shfl(d, threadIdx.x + 16);
    d += __shfl(d, threadIdx.x + 8);
    d += __shfl(d, threadIdx.x + 4);
    d += __shfl(d, threadIdx.x + 2);
    d += __shfl(d, 1);
}
```

Performance 12M Reduction


	Time ($3 \cdot 2^{22}$ floats)	Bandwidth	Step speedup	Cumulative Speedup	Speedup to reduction
AtomicGlobal	37.53ms	1.248GB/s	1.000	1.000	
AtomicShared	29.94ms	1.565GB/s	1.254	1.254	
Reduction	4.452ms	10.52GB/s	6.725	8.43	1.000
AntiDivergence	3.604ms	13.00GB/s	1.235	10.41	1.235
AntiConflicts	2.751ms	17.03GB/s	1.310	13.64	1.618
DualLoad	1.411ms	33.20GB/s	1.949	26.59	3.154
UnrolledLast	0.978ms	47.90GB/s	1.443	38.35	4.550
UnrolledAll	0.821ms	57.02GB/s	1.191	45.66	5.417
Shuffle	0.761ms	61.58GB/s	1.080	49.31	5.850

GTX 680 peak: 192GB/S

Reduction #9: SharedShuffle

- `__shfl(float val, int src)`

it is possible to directly exchange registers within a warp (CC \geq 3.0)



Reduction Complexity

- $\log(N)$ parallel steps
Each step S does $\frac{N}{2^S}$ independent ops
→ Step complexity is $O(\log N)$
- $N = 2^D$ operations = $\sum_{S \in [1..D]} 2^{D-S} = N - 1$
→ Work complexity is $O(N)$ → work efficient
→ Performs as many steps as sequential algorithm
- P threads physically in parallel (P processors)
→ Time complexity is $O\left(\frac{N}{P} + \log P\right)$
→ For a block $N = P$, so $O(\log N)$

Reduction Cost

- Cost of a parallel algorithm is processors \times time complexity
 - With N threads
 - Cost complexity = $N \cdot O(\log N) = \mathbf{O(N \log N)}$
 - not cost efficient
- Brent's theorem suggests $N/\log N$ threads
 - Each thread does $O(\log N)$ sequential work
 - Then all $N/\log N$ threads cooperate for $O(\log N)$ steps
 - Cost = $N/\log N \cdot O(\log N) + N/\log N \cdot O(\log N) = \mathbf{O(N)}$
 - Balance between sequential and parallel work
- For limited number of physical processors we see the same effect

Reduction #10: SharedMultiSeq

```
extern __shared__ float data[];
int id = threadIdx.x + blockIdx.x*blockDim.x*2;
data[threadIdx.x] = input[id] + input[id+blockDim.x];
__syncthreads();
```



```
extern __shared__ float data[];
int id = threadIdx.x + blockIdx.x*blockDim.x;
int gridSize = BlockSize*gridDim.x;

float in = 0.0f;
for(; id < elements; id += gridSize)
    in += input[id];

data[threadIdx.x] = in;
__syncthreads();
```


Performance 12M Reduction

	Time ($3 \cdot 2^{22}$ floats)	Bandwidth	Step speedup	Cumulative Speedup	Speedup to reduction
AtomicGlobal	37.53ms	1.248GB/s	1.000	1.000	
AtomicShared	29.94ms	1.565GB/s	1.254	1.254	
Reduction	4.452ms	10.52GB/s	6.725	8.43	1.000
AntiDivergence	3.604ms	13.00GB/s	1.235	10.41	1.235
AntiConflicts	2.751ms	17.03GB/s	1.310	13.64	1.618
DualLoad	1.411ms	33.20GB/s	1.949	26.59	3.154
UnrolledLast	0.978ms	47.90GB/s	1.443	38.35	4.550
UnrolledAll	0.821ms	57.02GB/s	1.191	45.66	5.417
Shuffle	0.761ms	61.58GB/s	1.080	49.31	5.850
MultiSeq	0.389ms	120.4GB/s	1.956	96.47	11.44


GTX 680 peak: 192GB/S

Reduction #10: SharedMultiSeq

SharedShuffle<<<12288,512>>>(...)



SharedMultiSeq<<<96,512>>>(...)



Questions

